

Given a corpus
no queries, no relevance labels

an LLM writes realistic queries
100 per task, from sampled documents

each model retrieves, a judge grades
every result 0–5 · or compares two models' top 10

then ranks the models
same top 3 as MTEB · ρ = 0.95