

Given a corpus
no queries, no relevance labels

an LLM writes realistic queries
DeepSeek-V4-Flash, 100 per task

each model retrieves, a decision model judges
28 embedding models · JEV-27B grades a result or picks a top 10

then ranks the models
same top 3 as MTEB · ρ = 0.95